

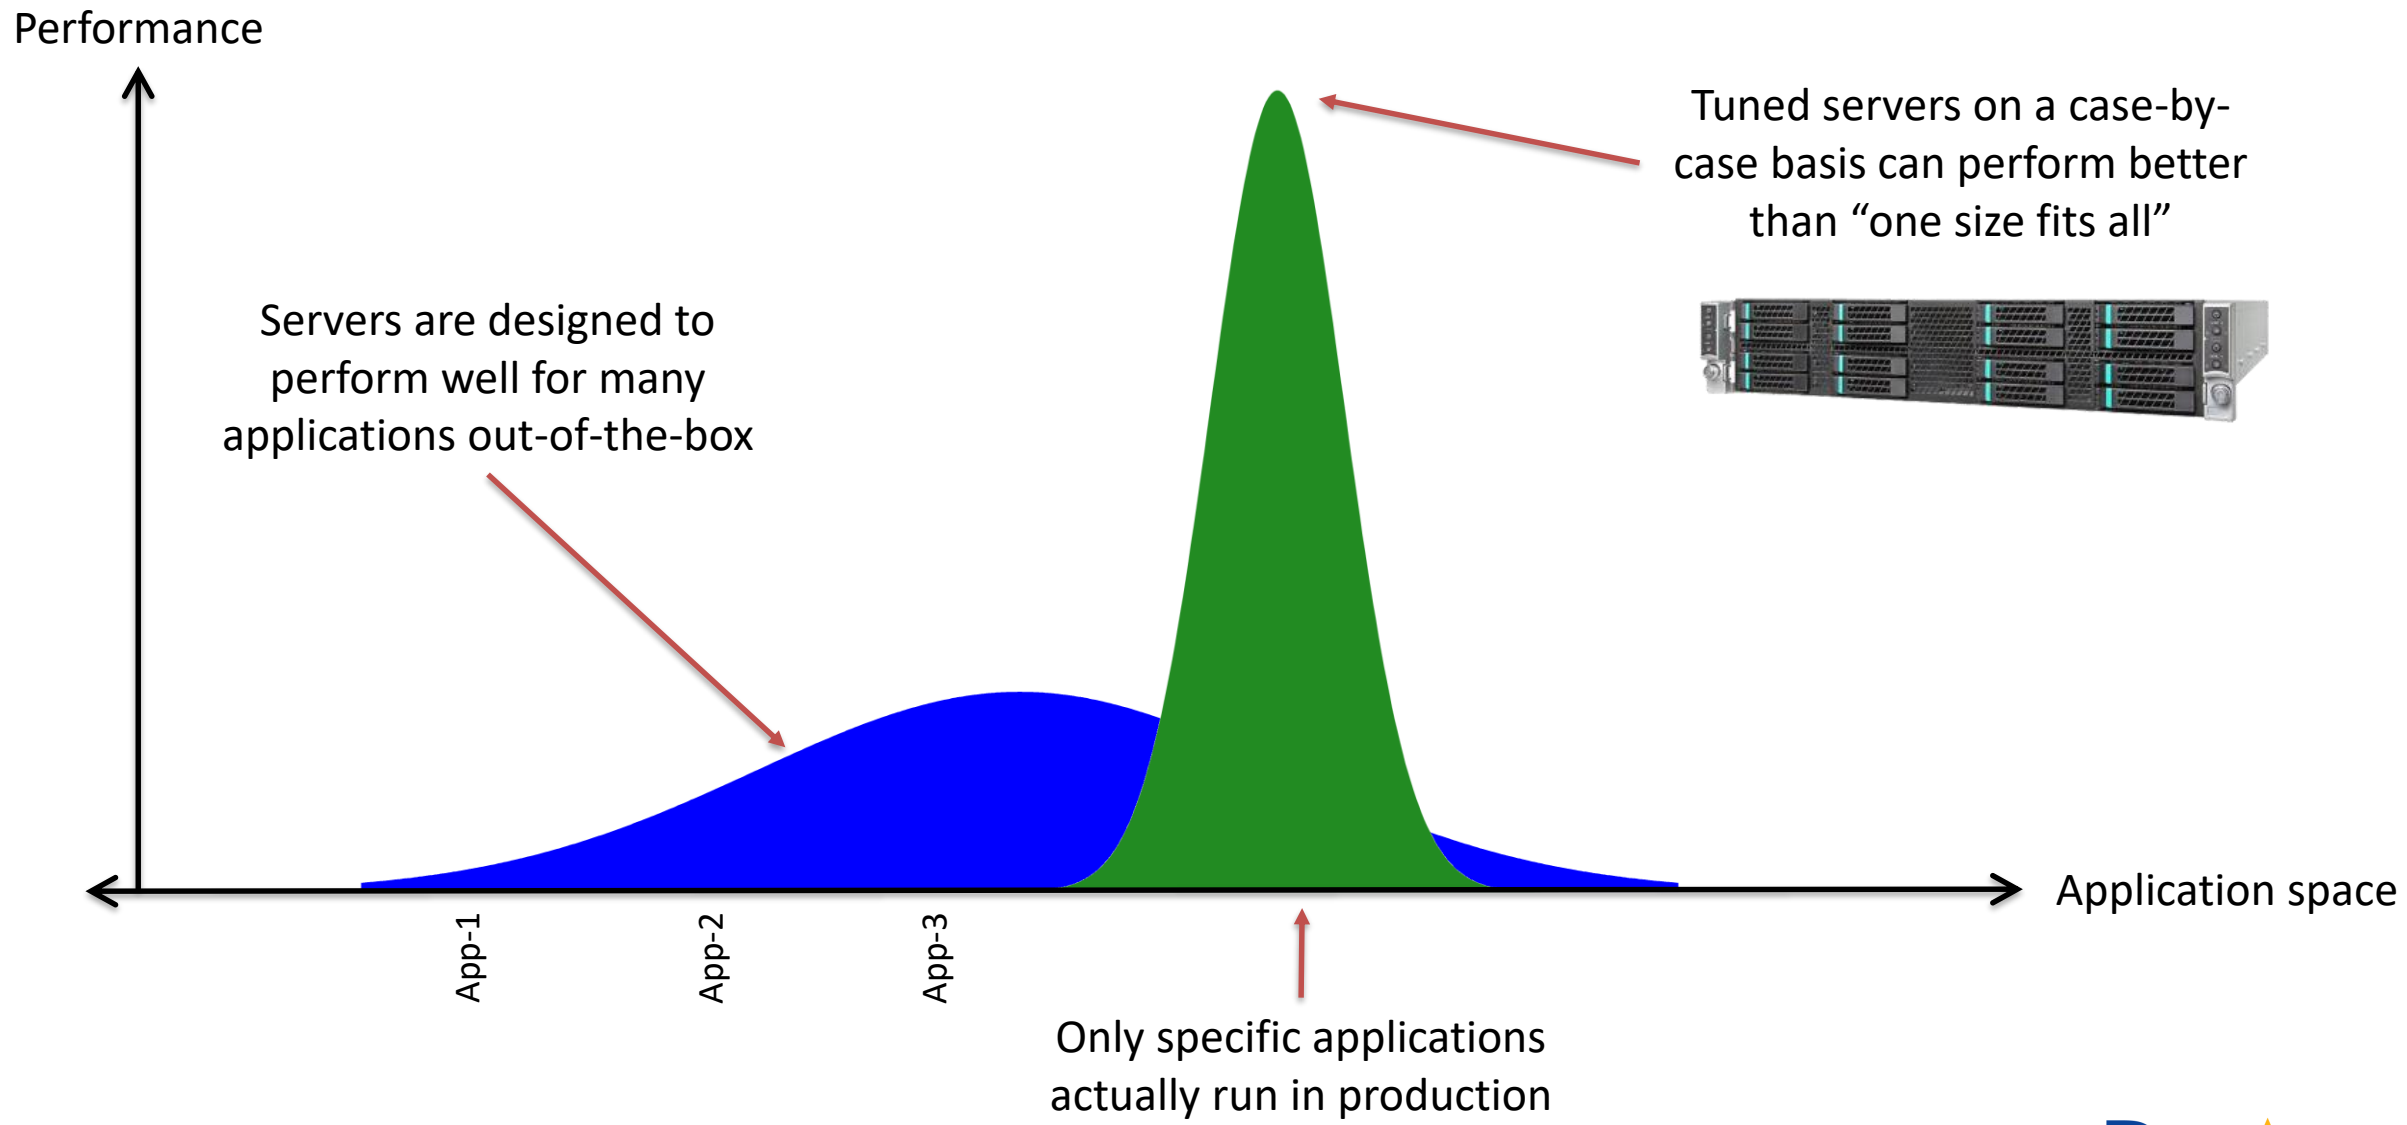
# The Era of Self-Tuning Servers

February 7, 2017



Dr. Tomer Morad, CEO  
[tomor.morad@datarcs.com](mailto:tomor.morad@datarcs.com)  
[www.datarcs.com](http://www.datarcs.com)

# Intuition



# Introduction to Tuning

❑ Knobs (in this talk) represent settings on a server that:




- ❑ Can be changed in real time
- ❑ Affects performance / energy efficiency
- ❑ Retains correctness

❑ Tuning

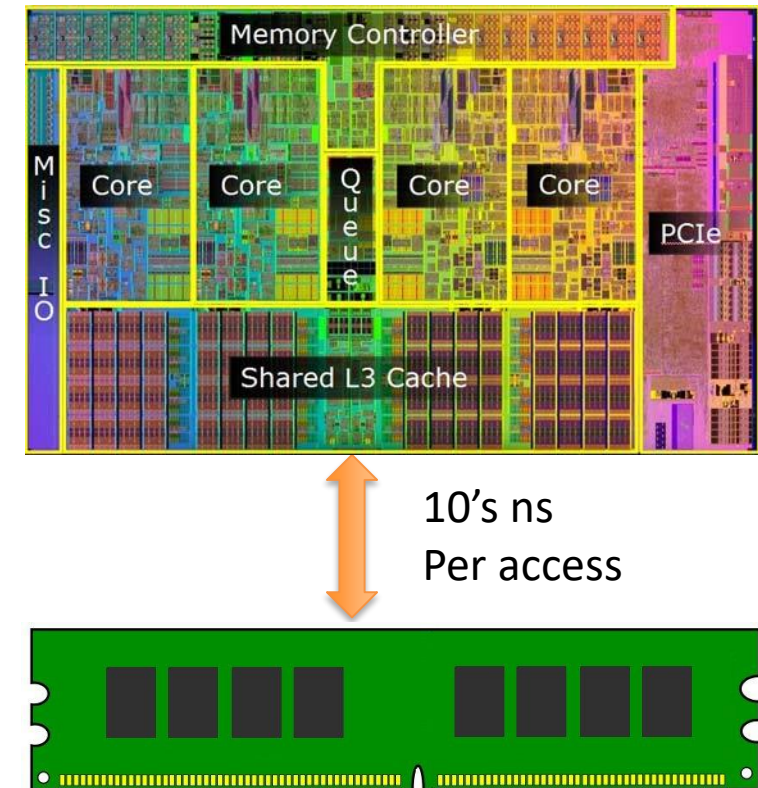


- ❑ The process of finding the best setting of a knob
- ❑ Tuning example (one application, one on/off knob):



# Example knob: CPU Cache Prefetching

- ❑ On-chip memory is 10x-100x faster than off-chip memory
- ❑ CPU Cache Prefetching – fetches data from off-chip memory **before** the CPU asks for it.
- ❑ Prefetcher for latency reduction:
  - ❑ **Predicts** which data the CPU will need in the future
  - ❑ **Predicts** which data the CPU will not require in the future
- ❑ Fetches data that the prefetcher predicted the CPU will need and store it in the cache instead of data that the prefetcher predicted the CPU will not need



```
[root@datarcs-ams-type1 ~]# wrmsr --all 0x1a4 0
[root@datarcs-ams-type1 ~]# ./demobench sequential
elapsed time: 6.352662865 seconds
[root@datarcs-ams-type1 ~]# wrmsr --all 0x1a4 7
[root@datarcs-ams-type1 ~]# ./demobench sequential
elapsed time: 9.297951911 seconds
[root@datarcs-ams-type1 ~]# wrmsr --all 0x1a4 0
[root@datarcs-ams-type1 ~]#
[root@datarcs-ams-type1 ~]#
[root@datarcs-ams-type1 ~]# ./demobench antiprefetch
elapsed time: 4.067213291 seconds
[root@datarcs-ams-type1 ~]# wrmsr --all 0x1a4 7
[root@datarcs-ams-type1 ~]# ./demobench antiprefetch
elapsed time: 3.501621960 seconds
[root@datarcs-ams-type1 ~]# wrmsr --all 0x1a4 0
```

- ❑ Programs have phases, for example:



- ❑ In this example:

```
[root@datarcs-ams-type1 ~]# wrmsr --all 0x1a4 0
[root@datarcs-ams-type1 ~]# ./demobench phases
elapsed time: 55.852871715 seconds
[root@datarcs-ams-type1 ~]# wrmsr --all 0x1a4 7
[root@datarcs-ams-type1 ~]# ./demobench phases
elapsed time: 56.189048409 seconds
[root@datarcs-ams-type1 ~]# wrmsr --all 0x1a4 0
```

- ❑ Can we do better?

# Tunable Knobs in Today's Systems

## Hardware

- Hardware prefetching
- SMT
- Cache partitioning
- Peripheral power states
- ...

## Firmware

- Power Management Unit (PMU), DVFS, Power States
- CPU Microcode
- ...

## Operating System

- Task Scheduler
- IO Scheduler
- Page Cache
- File Prefetching Algorithm
- Memory Allocation Algorithm
- Affinity
- ...

## Application

- Application-defined
- Choice of compiler
- Choice of libraries
- ...

**100's of different knobs to tune!**

# Some of the knobs (module-related) that can be tuned in Linux...

8250.nr_uarts	8250.share_irqs	8250.skip_txen_test	acpi.aml_debug_output	acpi.ec_busy_polling	acpi.ec_delay	acpi.ec_event_clearing	acpi.ec_polling_guard	acpi.ec_storm_threshold	acpi.immediate_undock	ahci.marvell_enable	apparmor.audit	apparmor.audit_header	apparmor.debug	apparmor.lock_policy	apparmor.logsyscall	apparmor.mode	apparmor.paranoid_load
apparmor.path_max	battery.cache_time	cfg80211.cfg80211_disable_40mhz_24ghz	debug_core.kgdbreboot	debug_core.kgdb_use_console	dm_mod.dm_mod_nr_hw_queues	dm_mod.dm_mod_queue_depth	dm_mod.dm_mod_numa_node	dm_mod.reserved_bio_based_ios	dm_mod.reserved_rq_based_ios	dm_mod.use_blk_mq	drm_kms_helper.dp_aux_i2c_speed_khz	drm_kms_helper.dp_aux_i2c_transfer_size	drm_kms_helper.edid_firmware	drm_kms_helper.fbdev_emulation	drm_kms_helper.poll	drm.debug	drm.timestamp_monotonic
drm.timestamp_precision_usecs	drm.vblankoff_delay	dynamic_debug.verbose	efi_pstore.pstore_disable	firmware_class.path	fuse.max_user_bgreq	fuse.max_user_congthresh	hid.debug	hid.ignore_special_drivers	i8042.debug	i8042.unmask_kbd_data	i915.enable_cmd_parser	i915.enable_fb_c	i915.enable_hangcheck	i915.enable_ips	i915.enable_psr	i915.fastboot	i915.invert_brightness
i915.load_detect_test	i915.lvds_use_ssc	i915.mmio_debug	i915.nuclear_pageflip	i915.panel_ignore_lid	i915.prefault_disable	i915.reset	i915.use_mmio_flip	i915.verbose_state_checks	ima.ahash_buffer_size	ima.ahash_min_size	intel_powerclamp.duration	intel_powerclamp.window_size	kdb.cmd_enable	kdb.enable_nmi	kernel.ignore_rlimit_data	kernel.initcall_debug	kernel.panic
kernel.panic_on_warn	kernel.pause_on_oops	keyboard.brl_n_bchords	keyboard.brl_timeout	kgdb_nmi.knocking	kgdb_nmi.magic	kgdboc.kgdboc	kvm.allow_unsafe_assigned_interrupts	kvm.halt_poll_ns	kvm.halt_poll_ns_grow	kvm.halt_poll_ns_shrink	kvm.ignore_msrs	kvm.lapic_timer_advance_ns	kvm.min_timer_period_us	kvm.tsc_tolerance_ppm	libahci.devslp_idle_timeout	libata.acpi_gtf_filter	libata.ignore_hpa
libata.zpodd_poweroff_delay	mac80211.beacon_loss_count	mac80211.ieee80211_default_rc_algo	mac80211.max_nullfunc_tries	mac80211.max_probe_tries	mac80211.min_strel_vht_only	mac80211.probe_wait_ms	module.sig_enforce	mousedev.tap_time	mousedev.xres	mousedev.yres	netpoll.carrier_timeout	nf_conntrack_ipv4.hashsize	nf_conntrack.acct	nf_conntrack.hashsize	nf_conntrack.nf_conntrack_helper	nf_conntrack.timestamp	overlay.check_copy_up
parport_pc.verbose_probing	pcie_aspm.policy	pciehp.pciehp_debug	pciehp.pciehp_force	pciehp.pciehp_poll_mode	pciehp.pciehp_poll_time	pci_hotplug.debug	pci_hotplug.debug_acpi	pci_slot.debug	ppp_generic.mpp_protocol_compress	printk.always_kmsg_dump	printk.console_suspend	printk.ignore_level	printk.time	processor.ignore_ppc	processor.ignore_tpc	processor.lateness_factor	psmouse.proto
psmouse.rate	psmouse.reset_after	psmouse.resolution	psmouse.resync_time	psmouse.smartscroll	pstore.update_ms	rcupdate.rcu_cpu_stall_suppress	rcupdate.rcu_cpu_stall_timeout	rcutree.jiffies_till_first_fqs	rcutree.jiffies_till_next_fqs	rcutree.jiffies_till_sched_qs	rcutree.kthread_prio	rng_core.current_quality	rng_core.default_quality	scsi_mod.default_dev_flags	scsi_mod.eh_deadline	scsi_mod.inq_timeout	scsi_mod.max_luns
scsi_mod.scsi_logging_level	scsi_mod.use_blk_mq	sg.allow_dio	sg.def_reserve_d_size	sg.scatter_element_size	shpchp.shpchp_debug	shpchp.shpchp_poll_mode	shpchp.shpchp_poll_time	snd_hda_codec_hdmi.static_pcm	snd_hda_codec.dump_coef	snd_hda_intel.align_buffer_size	snd_hda_intel.bdl_pos_adj	snd_hda_intel.power_save	snd_hda_intel.power_save_controller	snd_seq_midi.input_buffer_size	snd_seq_midi.output_buffer_size	snd_seq.seq_default_timer_card	snd_seq.seq_default_timer_class
snd_seq.seq_default_timer_device	snd_seq.seq_default_timer_resolution	snd_seq.seq_default_timer_class	snd_seq.seq_default_timer_subdevice	spurious_irqfixup	spurious_noirq_debug	sr_mod.xa_test	suspend.pmt_est_delay	sysrq.reset_seq	sysrq.sysrq_downdtime_ms	tcp_cubic.beta	tcp_cubic.fast_convergence	tcp_cubic.hystart	tcp_cubic.hystart_ack_delta	tcp_cubic.hystart_detect	tcp_cubic.hystart_low_window	tcp_cubic.initial_ssthresh	tcp_cubic.tcp_friendliness
thermal.act	thermal.crt	thermal.psv	tpm.suspend_pcr	usbcore.authorized_default	usbcore.autosuspend	usbcore.initial_descriptor_timeout	usbcore.old_scheme_first	usbcore.usbfs_memory_mb	usbcore.usbfs_snoop	usbcore.usbfs_snoop_max	usbcore.use_both_schemes	usbhid.ignoreled	usbhid.mousepoll	usb_storage.delay_use	usb_storage.option_zero_cd	usb_storage quirks	usb_storage.sanitize_install
video.allow_duplicates	video.brightness_switch_enabled	video.report_key_events	vt.color	vt.cur_default	vt.default_blue	vt.default_green	vt.default_red	vt.default_utf8	vt.global_cursor_default	vt.italic	vt.underline	workqueue.debug_force_rr_cpu	workqueue.watchdog_thresh	x86_pkg_temp_thermal.notification_delay_ms	xhci_hcd.link_quirk	zswap.compressor	zswap.enabled
								zswap.max_pool_percent	zswap.zpool								



# Limitations of manual tuning

100's of knobs –  
too many to  
manually tune

Dependencies  
among different  
knobs

Knob settings  
depend on  
hardware

Knob settings  
depend on  
applications and  
input data

No practical way to  
“see” program  
phases

Labor intensive  
task

Requires expertise  
which is absent in  
most organizations

# The Era of Self-Tuning Servers - DatArcs Optimizer©



Step 1:  
Install

Step 2:  
Learn

Step 3:  
Optimize



# DatArcs Optimizer Advantages

## Automatic

Optimizer tunes without  
any user input

## Adaptive

Program phase and  
workload detection

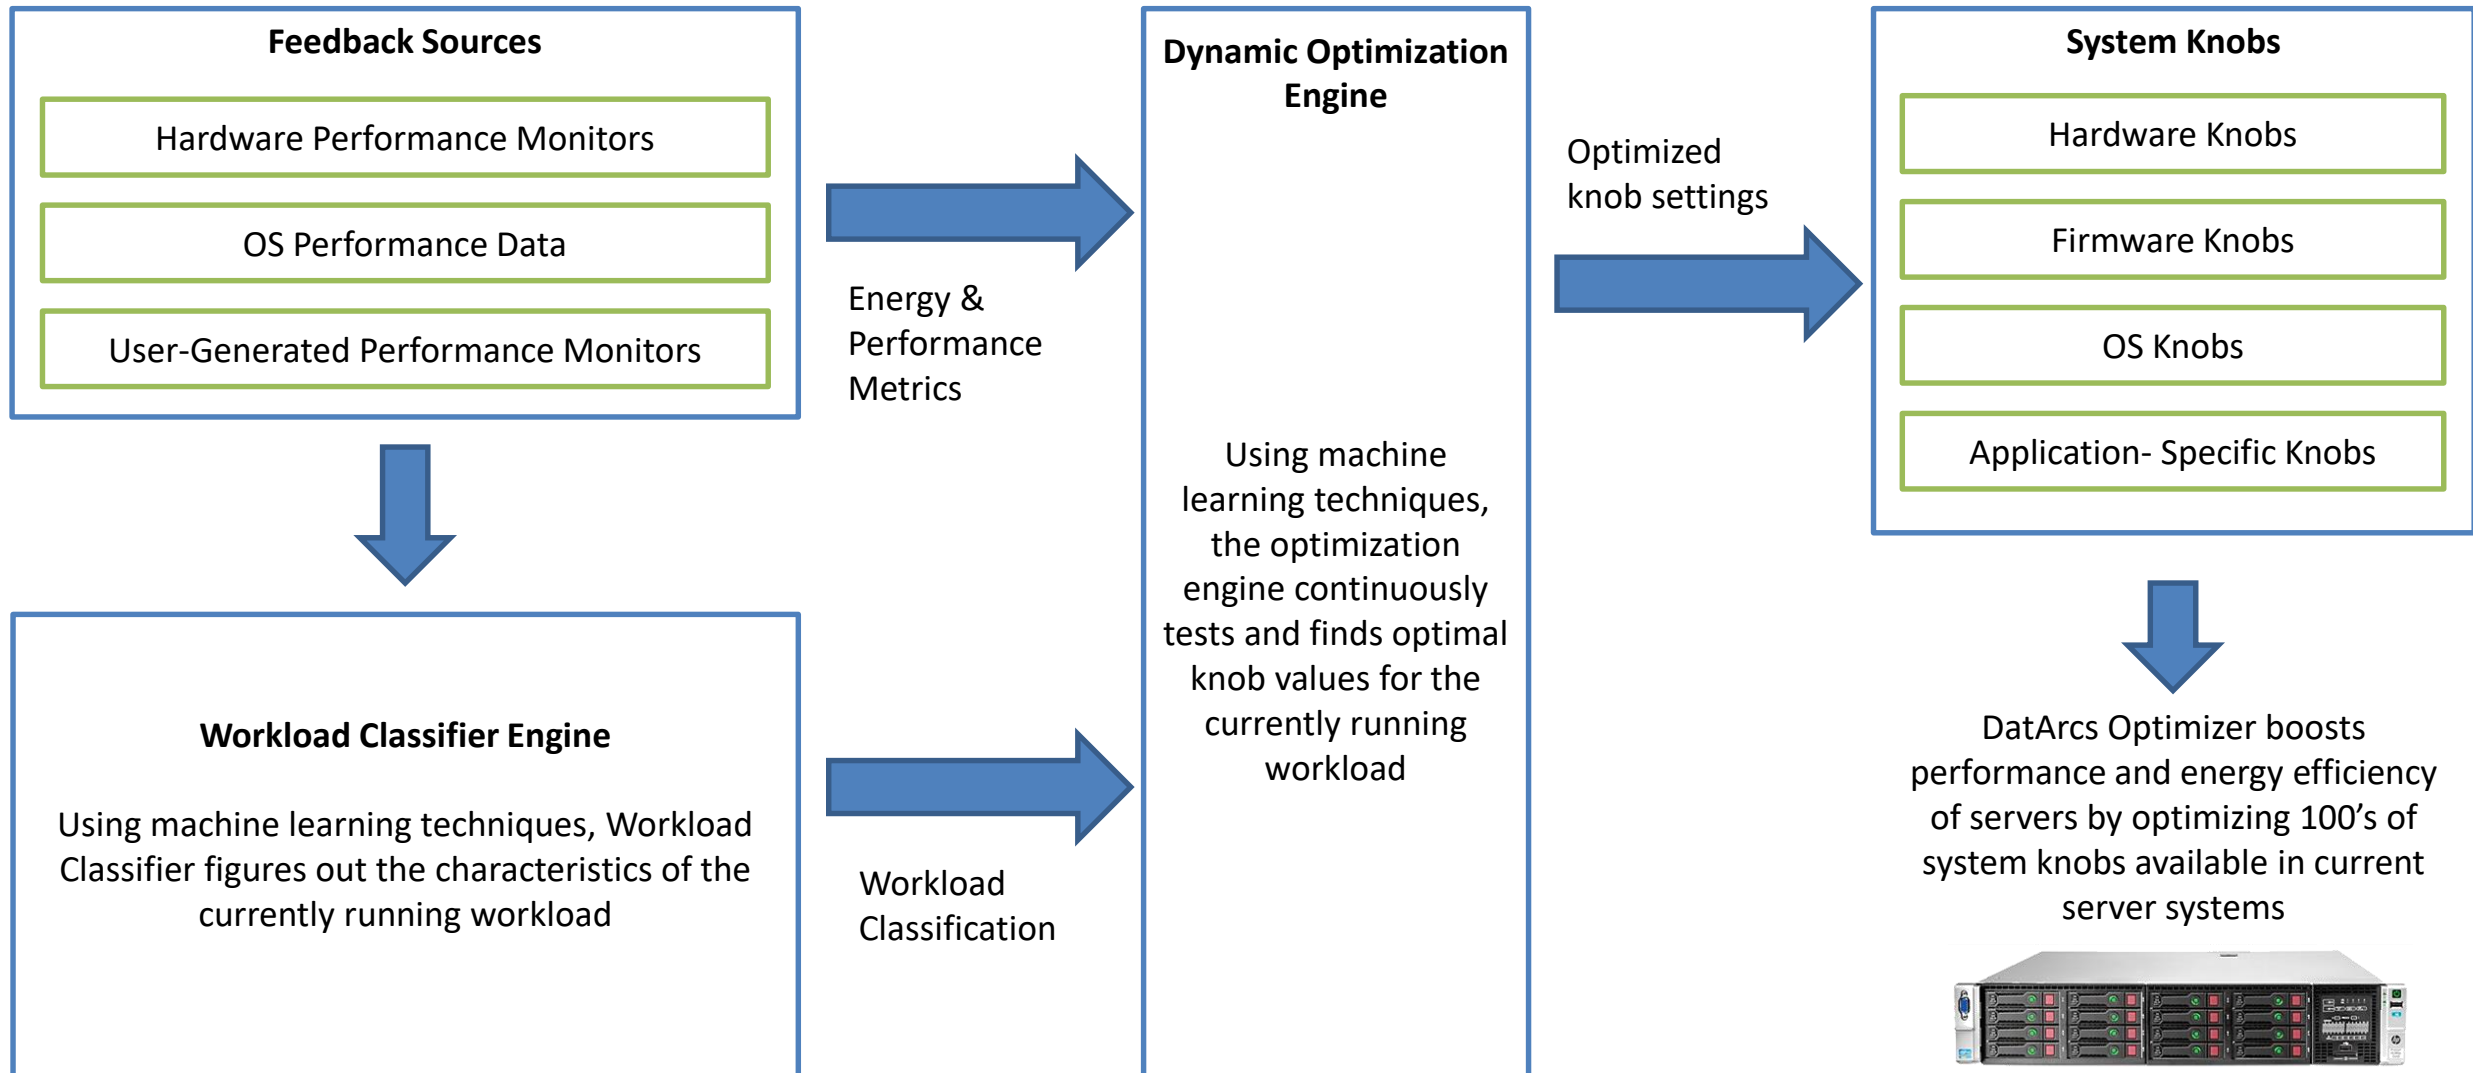
## Flexible

Tunes for speed, energy,  
power, power cap, etc.

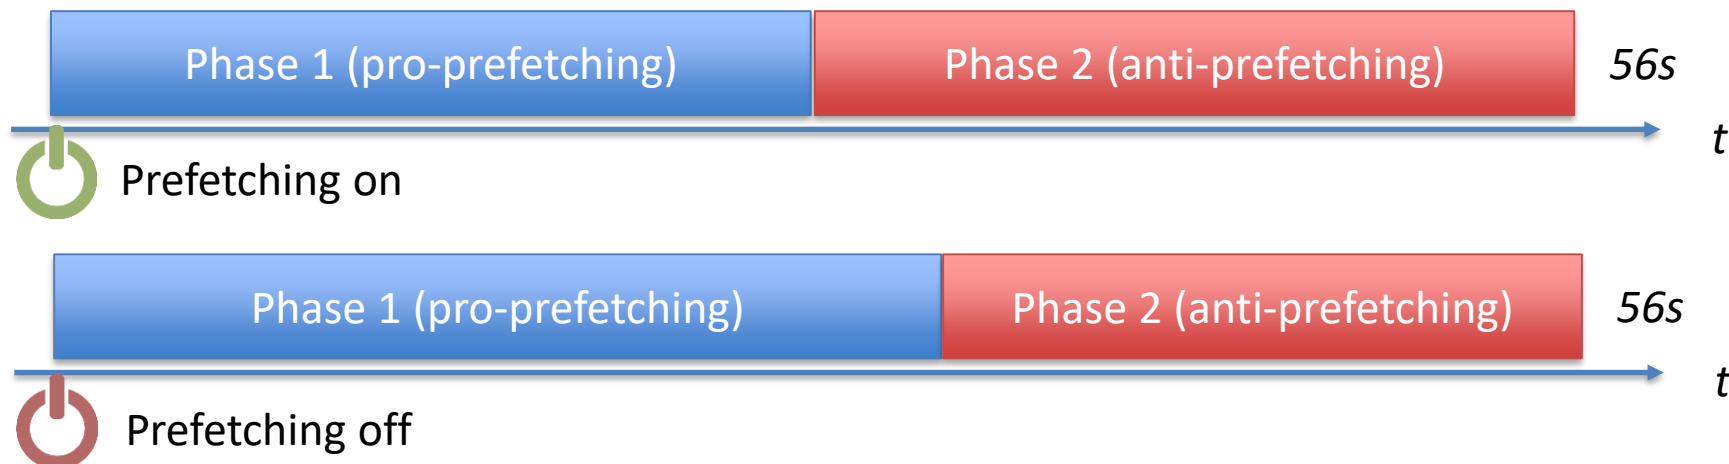
## Extensible

Users can add application-  
specific knobs

# DatArcs Optimizer In a Nutshell




## Manual Tuning



## Dynamic Tuning


```
# systemctl start datarcs-optimizer
# ./demobench phases -1
elapsed time: 56.668542269 seconds
elapsed time: 53.045306135 seconds
elapsed time: 52.650779192 seconds
elapsed time: 55.153519531 seconds
elapsed time: 51.656128995 seconds
elapsed time: 52.576210238 seconds
elapsed time: 51.734081272 seconds
elapsed time: 51.160347630 seconds
elapsed time: 53.263679141 seconds
elapsed time: 51.071073557 seconds
elapsed time: 51.354038358 seconds
```



**DatArcs Optimizer outperformed the manual method!**



### DatArcs Optimizer Performance



#### Benchmark

- Phoronix test of Apache web server

#### Experiment phases

- **Baseline** - 20 runs without Optimizer
- **Learning** - 140 runs with Optimizer in learning mode
- **Optimization** - 20 runs with Optimizer in optimization mode
- **Static** - 20 runs after Optimizer applied the best knob values and exited

#### Server Configuration

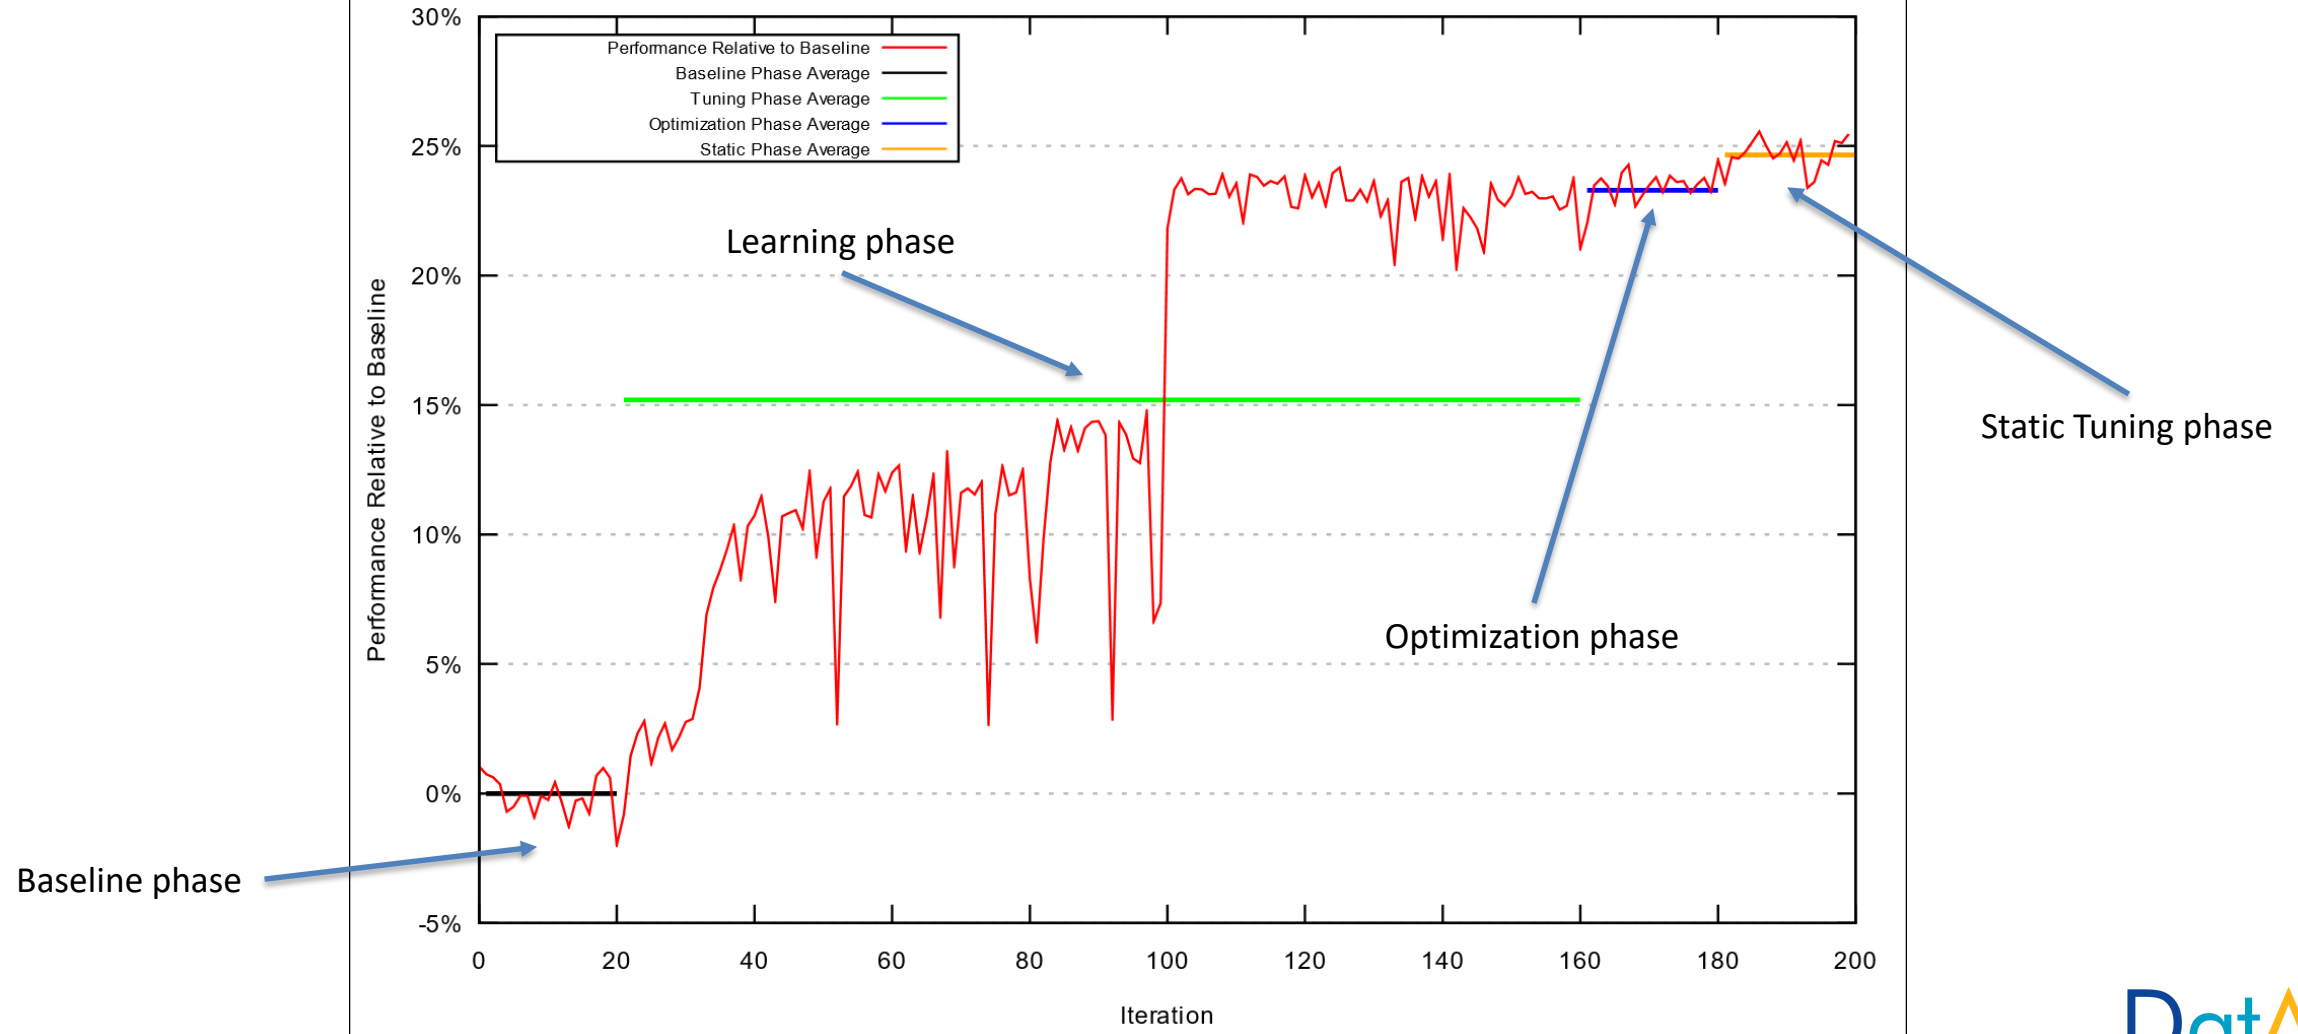
- Packet type-2 server: 2x Intel E5-2650 @ 2.2GHz, 24 cores, 48 threads

#### Results

- **DatArcs Optimizer correctly identified the effects of various knobs on the application performance, and achieved ~25% boost in performance**



### DatArcs Optimizer Performance



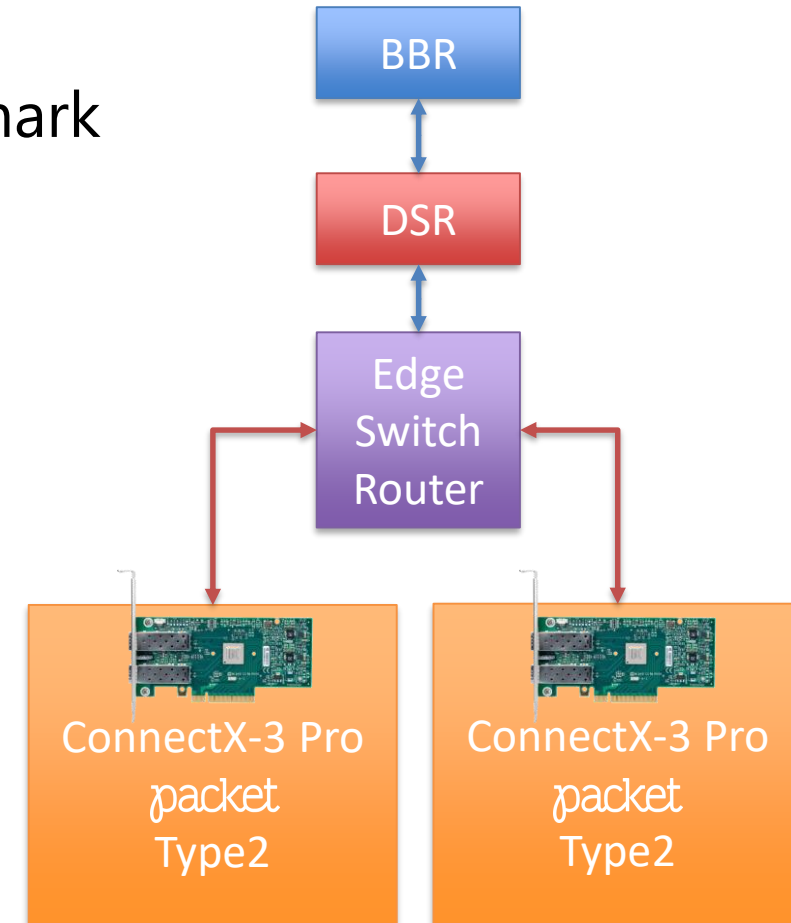
- ❑ One knob: tx-usecs = (16 or 256)
- ❑ Two benchmarks, each prefers a different setting
- ❑ DatArcs Optimizer detects best setting for each benchmark

```
# ./demo_mellanox 1
MIGRATED TCP STREAM TEST from 0.0.0.0
(0.0.0.0) port 0 AF_INET to 147.75.108.33
() port 0 AF_INET : demo
Interim result: 4496.55 10^6bits/s
Interim result: 4504.32 10^6bits/s
Interim result: 4556.30 10^6bits/s
Interim result: 4511.66 10^6bits/s
Interim result: 4490.83 10^6bits/s
Interim result: 8089.87 10^6bits/s
Interim result: 9259.31 10^6bits/s
Interim result: 9276.94 10^6bits/s
Interim result: 9258.05 10^6bits/s
Interim result: 9263.36 10^6bits/s
Interim result: 9218.19 10^6bits/s
Interim result: 9260.11 10^6bits/s
Interim result: 9255.66 10^6bits/s
Interim result: 9265.10 10^6bits/s
Interim result: 9264.68 10^6bits/s
Interim result: 9392.09 10^6bits/s
Interim result: 9267.17 10^6bits/s
```

**+106%**

```
# ./demo_mellanox 2
MIGRATED TCP STREAM TEST from 0.0.0.0
(0.0.0.0) port 0 AF_INET to 147.75.108.33
() port 0 AF_INET : demo
Interim result: 1000.47 10^6bits/s
Interim result: 988.52 10^6bits/s
Interim result: 1012.41 10^6bits/s
Interim result: 1104.88 10^6bits/s
Interim result: 996.59 10^6bits/s
Interim result: 1047.43 10^6bits/s
Interim result: 1815.29 10^6bits/s
Interim result: 1600.28 10^6bits/s
Interim result: 1890.25 10^6bits/s
Interim result: 2052.70 10^6bits/s
Interim result: 2086.82 10^6bits/s
Interim result: 1973.90 10^6bits/s
Interim result: 2038.70 10^6bits/s
Interim result: 2026.54 10^6bits/s
Interim result: 2003.32 10^6bits/s
Interim result: 2062.60 10^6bits/s
Interim result: 2097.10 10^6bits/s
```

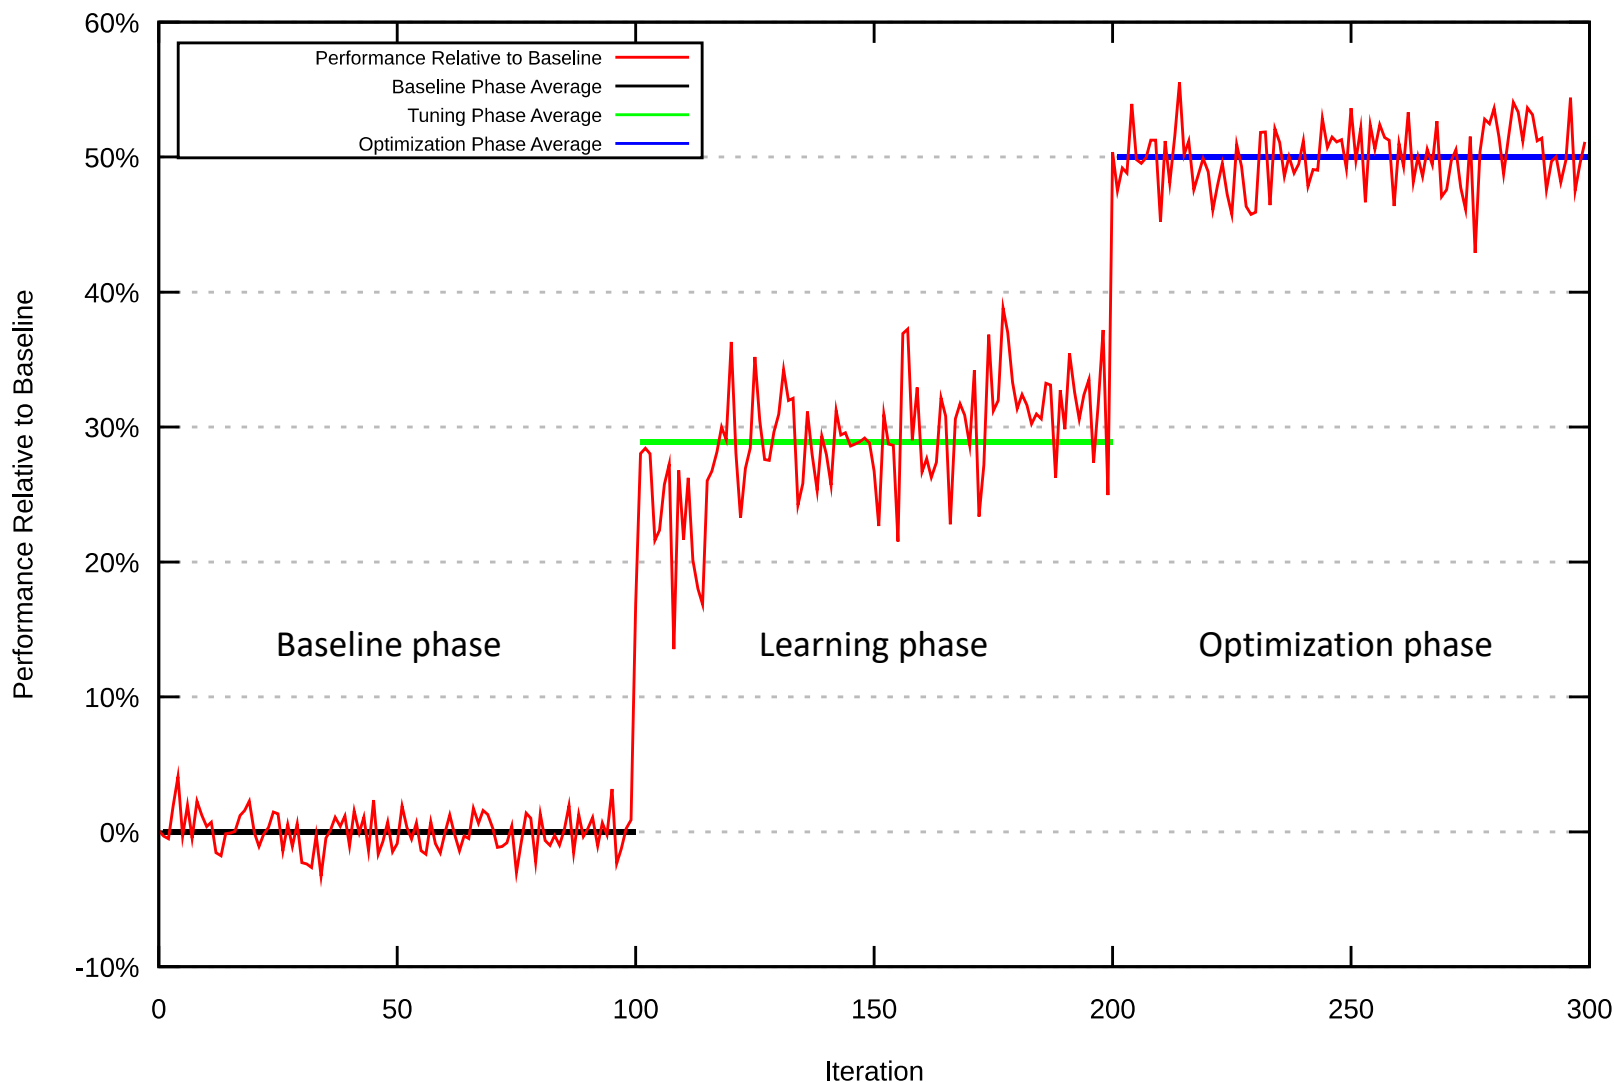
**+109%**







### DatArcs Optimizer Performance



#### Benchmark

- Collaborative Filtering using Apache Spark
- Three containers on a single machine:
  - One Master
  - Two Workers

#### Server Configuration

- Packet type-2 server: 2x Intel E5-2650 @ 2.2GHz, 24 cores, 48 threads

# Dynamic Tuning Trends

## Number of knobs is rising

- Hardware becomes more complex and configurable
- Operating systems also become more complex and configurable
- Heterogeneity: GPUs, Accelerators, FPGAs, etc.



## Expertise is becoming scarce

- Wrong to assume that tuning is being taken care of by the “cloud”



## Hardware matters!

- More tuning opportunities
- VMs Vs. Bare Metal and Containers



# Summary

- ✓ **Server tuning as we know it is about to change!** Dynamic tuning is a quick and simple way to improve performance and energy efficiency of existing systems
- ✓ DatArcs Optimizer is the first feedback-based dynamic tuning software
- ✓ Significant improvement in benchmarks (~50%), and we've only just begun
- ✓ Version 0.5 now in closed beta – feedback welcomed!

DatArcs Optimizer Performance

